

---

# Reinforcement Learning TP1 : Dynamic Programming and Reinforcement Learning

---

Plassier Vincent

## 1 Dynamic Programming

### 1.1

We can guess that  $[1, 1, 2]$  is the optimal policy.

### 1.2

Thanks to the Puterman's theorem with  $\varepsilon = \frac{1-\gamma}{200\gamma}$ , once can deduce that :

$$\|v^{k+1} - v^k\|_{\infty} \implies \|v^{\pi_{k+1}} - v^{\star}\|_{\infty} \leq 0.01$$




FIGURE 1:  $\|v^{\pi_k} - v^{\star}\|_{\infty}$  in function of k

### 1.3

Now, we display the iterations obtained by the algorithm Policy iteration (PI) :




FIGURE 2:  $\|v^{\pi_k} - v^*\|_{\infty}$  in function of k

We find that the Political Iteration (PI) algorithm converges in only 3 iterations whereas the algorithm Value iteration (VI) need more run. However, the PI algorithm solves a linear system. So PI has a complexity about  $\mathcal{O}(n_{states}^3)$ , whereas VI has a complexity about  $\mathcal{O}(n_{action}^2 \times n_{states})$ . Thus, it may be advantageous to choose the method based on the ratio  $n_{action}/n_{states}$ .

## 2 Reinforcement Learning

### 2.1

In this question, we consider the quantity  $\sum_{s \in S} \mu_0(s) [V_n(s) - V^{\pi}(s)]$  above  $n$ . We plot the first iterations :




FIGURE 3:  $\sum_{s \in S} \mu_0(s) [V_n(s) - V^{\pi}(s)]$  in function of  $n$

The error this decay to 0 when  $n$  increase.

### 2.2

From now on, we are interested by the study of a Q-learning algorithm. In the algorthime described in the subject, we can change the exploration-exploitation's rapport by playing on the parameters  $\varepsilon$  and  $\alpha_i(x, a)$ . The first figure is obtained by considering  $\varepsilon_t = (0.99)^t$ . Thus, as the number of iterations increase, our learning rate decreases. Therefore, the algorithm will more and more privilege exploitation.




FIGURE 4: Error and reward evolutions in function of the epoch

We also represent the evolution of  $\|v^* - v_{\pi_n}\|_\infty$  during learning as well as the sum of cumulative rewards. Considering this time  $\varepsilon \in \{0.01, 0.1\}$ . We get :




FIGURE 5: Representation des approximations obtenues

For the value  $\varepsilon = 0.01$ , we see that the algorithm still has not converged after 150 iterations. Indeed, there has not been enough exploration.

Remark : in order to determine the optimal hyperparameters, it is possible to use the python fmin function.

### 2.3

Let's denote  $V^*$  the unique fixed point of Gamma and  $\pi^*$  the optimal policy. According to the course,

$$\pi^*(x) = \arg \max_a \left[ r(x, a) + \sum_y p(y|x, a) V^*(y) \right]$$

The previous equation is independant of the choice of the initial policy  $\mu_0$ . Thus, we conclude that  $\pi^*$  is independant of  $\mu_0$ .

Numerically, we can verify that the choice of another initial distribution does not affect the result. For that, we draw the curve obtained by replacing the initial distribution :




FIGURE 6: Evolution of the error with an other initial distribution

We constate that the curve obtained [6](#) is very similar to the previous one in [5](#). We deduce that  $\mu_0$  input distribution does not influence the optimal policy.